

STANLEY CAO

caostan02@gmail.com | linkedin.com/in/stanley-cao | github.com/stanley-cao | stanleycao.vercel.app

Education

University Of Toronto

Expected June 2028

Masters of Science in Applied Computing, Artificial Intelligence

Toronto, Ontario

- **Coursework:** Neural Networks & Deep Learning, Machine Learning, Computational Imaging, Advanced Data Systems

Western University

October 2025

Honours Specialization in Computer Science

London, Ontario

Technical Skills

ML/AI: PyTorch, PyTorch Lightning, scikit-learn, YOLOv8, LLMs, RAG, pgvector, AI Agents, Claude Code, Cursor

Languages: Python, Java, C/C++, C#, JavaScript, TypeScript, SQL, HTML, CSS, R, Swift

Technologies/Frameworks: React, Next.js, Node.js, PostgreSQL, Docker, Git/GitHub, AWS, Supabase, FastAPI

Experience

Super.com

September 2026 – Present

Software Engineer Intern

Toronto, ON, Canada

- Closed a financial visibility gap by designing a scalable, partner-agnostic discount attribution framework, backfilling 1.3K+ historical bookings to trace \$8.5K in previously unattributed Hopper discounts
- Shipped a double-CTA credit feature to web and mobile booking confirmation pages on a platform serving 30M+ users, letting customers save past payment methods for faster repeat bookings
- Refactored the hotel booking flow to decouple 6 concluded A/B experiments and their stale feature flags from core production logic, simplifying checkout code paths and reducing regression risk

Morphace

October 2025 – December 2025

Software Engineer Intern

Toronto, ON, Canada

- Led development of a full-stack iOS mobile application using Swift and SwiftUI, building 6+ interactive features including product reviews, rating systems, and multi-select content filters to enhance user engagement and content discovery
- Implemented and integrated RESTful APIs (10+ endpoints) using FastAPI and Python, enabling end-to-end functionality for product discussions, discussion threads, posts, reactions, and ratings within the mobile app
- Architected a relational PostgreSQL database schema with 10+ tables, enforcing tag constraints and aggregated rating logic to ensure data consistency, scalability, and efficient backend operations

Wouessi Inc

February 2025 – April 2025

Software Engineer Intern

Toronto, ON, Canada

- Built and standardized 15+ reusable React components, including buttons, typography, and layout elements, enhancing design consistency, reducing UI duplication and accelerating feature deployment across the application
- Developed and optimized dynamic, filterable data tables in React.js and CSS Modules, improving data visualization speed, responsiveness and user interactivity by 25% under real world usage
- Performed front-end testing, debugging, and code reviews to resolve 10+ UI issues, improving cross-browser compatibility, and ensure stable, polished functionality prior to stakeholder demonstrations

Johnson & Johnson

May 2024 – August 2024

Data Engineer Intern

Toronto, ON, Canada

- Automated Fair Market Value Tableau dashboards by implementing 5+ reusable calculated fields, data transformations, and Excel based data pipelines, improving refresh efficiency and supporting real-time compliance analytics workflows
- Implemented and trained an AI image classification pipeline using YOLOv8 in Python, integrating automated preprocessing and inference logic to reduce manual review time by 40% across compliance workflows
- Enhanced compliance data categorization processes by assisting with SQL-based data analysis and validation workflows, improving tracking accuracy and enforcement reliability by 15% for healthcare professional interactions

Projects

RAG Assistant - Agentic Document Chat Platform | *Python, FastAPI, React, TypeScript, Supabase, pgvector*

- Built a full-stack RAG application supporting 10MB+ document processing through a 3-stage ingestion pipeline, with real-time progress tracking across 4 database tables for upload and processing state management
- Implemented an AI agent system with a 3-round decision loop enabling autonomous semantic search and retrieval using vector embeddings with configurable similarity thresholds
- Developed a multi-user platform with encrypted API key storage, row-level security across 5 tables, and real-time chat

EKGram – Deep Learning EKG Digitization | *PyTorch Lightning, Next.js, Supabase*

- Developing a PyTorch Lightning model with a 4-person team to convert scanned EKG printouts into time-series signals, training on 85 GB of 8,700+ images across 6 artifact types (damage, mold, grayscale, mobile photos)
- Leading model deployment through a cloud-hosted web app with live signal visualization and CSV export, and designing robustness tests for daily batch uploads, malformed inputs and out-of-distribution images